

gdbot — PPO Training Report


run "live" — conv policy over the 24x16x4 occupancy grid

PPO updates: 1,382 Environment steps: 1,412,608.0 Episodes: 1,760 Wall clock: 7.1 h

Mean progress — first update: 0.0% final: 25.4% best single attempt: 52.4%


Completions: 0 Throughput: 56 steps/s Final policy entropy: 0.238 (max 0.693)

Learning curve — how far the agent gets as it trains




Is the run healthy?


Mean episode return




Policy entropy — is it committing?




KL per update — step size



Value loss — can the critic predict returns?



Every attempt — the spread is what the mean above hides



Distribution of attempts — where the agent keeps dying

