

gdbot — PPO Training Report


run "simlong" — conv policy over the 24x16x4 occupancy grid

PPO updates: 400 Environment steps: 819,200.0 Episodes: 1,265 Wall clock: 0.1 h

Mean progress — first update: 16.2% final: 84.2% best single attempt: 99.9%


Completions: 455 Throughput: 1,860 steps/s Final policy entropy: 0.460 (max 0.693)

Learning curve — how far the agent gets as it trains




Is the run healthy?


Mean episode return




Policy entropy — is it committing?




KL per update — step size



Value loss — can the critic predict returns?



Every attempt — the spread is what the mean above hides



Distribution of attempts — where the agent keeps dying

