

gdbot — PPO Training Report


run "live-v2" — conv policy over the 24x16x4 occupancy grid

PPO updates: 211 Environment steps: 432,128.0 Episodes: 732 Wall clock: 2.2 h

Mean progress — first update: 3.0% final: 16.1% best single attempt: 19.6%


Completions: 0 Throughput: 54 steps/s Final policy entropy: 0.492 (max 0.693)

Learning curve — how far the agent gets as it trains




Is the run healthy?


Mean episode return




Policy entropy — is it committing?




KL per update — step size



Value loss — can the critic predict returns?



Every attempt — the spread is what the mean above hides



Distribution of attempts — where the agent keeps dying

